

```

#!/usr/bin/env python
# coding: utf-8

# # Question 1
#
# ## Import Data

# In[1]:

# Code to import "Alternative Fuel Vehicles US.csv"

import pandas as pd
data_1 = pd.read_csv('Alternative Fuel Vehicles US.csv')
data = pd.read_csv('Alternative Fuel Vehicles US.csv')

# ## Part 1 - a)

# In[2]:

# Filter dataframe for European and Japanese makes

Japanese = ['Hino', 'Hyundia', 'Toyota', 'Honda', 'Lexus']
European= ['Audi', 'BMW', 'Jaguar', 'Mercedes-Benze', 'Land Rover']
fuelmodel = data[['Manufacturer', 'Fuel', 'Conventional Fuel Economy Combined']]
fuelmode_final_European =
fuelmodel[(fuelmodel['Manufacturer'].isin(European))&(fuelmodel['Fuel']=='Hybrid
Electric')].dropna()
fuelmode_final_Japanese =
fuelmodel[(fuelmodel['Manufacturer'].isin(Japanese))&(fuelmodel['Fuel']=='Hybrid
Electric')].dropna()

# In[3]:

# Plot histograms side-by-side

import matplotlib.pyplot as plt
plt.figure()
plt.subplot(2,2,1)
plt.hist(fuelmode_final_European['Conventional Fuel Economy Combined'])
plt.title('Fuel economy of European')
plt.xlabel('Fuel economy')
plt.ylabel('Number')

plt.subplot(2,2,2)
plt.hist(fuelmode_final_Japanese['Conventional Fuel Economy Combined'])
plt.title('Fuel economy of Japanese')
plt.xlabel('Fuel economy')
plt.ylabel('Number')

# It is clear from the histogram that European cars are more fuel efficient than
Japanese cars

```

```

#

#

# ## Part 1 - b)

# In[4]:

# Code for plotting histogram
plt.figure()
plt.hist(fuelmode_final_Japanese['Conventional Fuel Economy Combined'])
plt.title('Fuel economy of Japanese Hybird Electric cars')
plt.xlabel('Fuel economy')
plt.ylabel('Number')
plt.show()

# In[5]:

# Code for getting and saving "TopTenElectric.csv"
All_Electric_Range = data[(data['Fuel'] == 'Electric')]
All_Electric_Range = All_Electric_Range.sort_values(by = ['All-Electric
Range'], ascending = False)
num= round(All_Electric_Range.shape[0]/10)
All_Electric_Range_final = All_Electric_Range[0:num]
All_Electric_Range_final
All_Electric_Range_final.to_csv("TopTenElectric.csv")

# ## Part 2 - a)

# Pick out the 'Economy Combined' entries for those vehicles that have a non-empty
'conventional fuel'.
# After setting column A, what are D, K, L, M and Q, fill in the NaNs in these
columns according to the mode or average.
# Then convert the feature columns except the target column to numerical variables
through one-hot encoding

# In[6]:

# Code for filtering dataset and other pre-processing applied

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import numpy as np
Economy_Combined = data[~data['Conventional Fuel Economy Combined'].isnull()]
A = 'Category'
D = 'Manufacturer'
K = 'Conventional Fuel Economy Combined'
L = 'Alternative Fuel Economy City'
M = 'Alternative Fuel Economy Highway'
Q = 'Drivetrain'
E = 'Fuel'
Euro_pean = ['Audi', 'Land Rover', 'BMW', 'Mercedes-
Benz', 'Jaguar', 'Volvo', 'Mini', 'Porsche', 'Bentley Motors', 'Ferrari']

```

```

American
=['Chevrolet','GMC','Cadillac','Ford','Jeep','Ram','Lincoln','Chrysler','Karma','Mitsubishi','Polestar Automotive USA']
Asian_manufacturers = ['Hyundai','Toyota','Honda','Kia','Lexus','Acura','Subaru']
cols = [A,D,K,L,M,Q,E]
dx = Economy_Combined[cols]
data = dx.copy()
data["Alternative Fuel Economy City"].fillna(dx["Alternative Fuel Economy City"].median(skipna=True), inplace=True)
data["Alternative Fuel Economy Highway"].fillna(dx["Alternative Fuel Economy Highway"].median(skipna=True), inplace=True)
data["Drivetrain"].fillna(dx["Drivetrain"].value_counts().idxmax(), inplace=True)
data['Fuel'] = data['Fuel'].astype('category')
data['Fuel'] = data['Fuel'].cat.codes
data_clean = data
data_clean = pd.get_dummies(data_clean[cols])

```

```
# ## Part 2 - b)
```

```
# In[7]:
```

```
# Get average accuracy of logistic regression across 20 random train/test splits
# You can manually change test_frac and present only the best value here
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
new_cols = ['Conventional Fuel Economy Combined','Alternative Fuel Economy City','Alternative Fuel Economy Highway','Category_Passenger Van/Shuttle Bus','Category_Pickup','Category_SUV','Category_Sedan/Wagon','Category_Van','Manufacturer_Acura','Manufacturer_Porsche','Manufacturer_Ram','Manufacturer_Subaru','Manufacturer_Toyota','Manufacturer_Volvo','Drivetrain_4WD','Drivetrain_AWD','Drivetrain_FWD','Drivetrain_Part-Time 4WD','Drivetrain_RWD']
j = 0
a_list = []
list_rate = [0.1,0.15,0.2,0.25,0.3]
while j <=4:
    a = 0
    i = 0
    while i <= 19:
        X = data_clean[new_cols]
        y = data_clean['Fuel']
        X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=list_rate[j])
        logreg = LogisticRegression()
        logreg.fit(X_train, y_train)
        y_pred = logreg.predict(X_test)
        a += accuracy_score(y_test, y_pred)
        i+=1
    a = a/20
    a_list.append(a)
    j+=1

print('Train/Test split results:')
print("The best accuracy : %2.3f" % max(a_list))

```

```
# ## Part 2 - c)
```

```
# In[8]:
```

```
# Average accuracy of MLP across 20 random train/test splits
# You can manually change number of neurons in hidden layer and present best result
from sklearn.neural_network import MLPClassifier
i = 0
acc_final = 0
while i<= 19:
    df_train,df_test = train_test_split(data_clean, test_size=0.3)
    X_train = df_train[new_cols].to_numpy()
    y_train = df_train['Fuel'].to_numpy()

    X_test = df_test[new_cols].to_numpy()
    y_test = df_test['Fuel'].to_numpy()

    clf = MLPClassifier(random_state = 1,hidden_layer_sizes =
(100, )).fit(X_train,y_train)

    y_pred = clf.predict(X_test)

    acc = accuracy_score(y_test, y_pred)
    acc_final = acc_final + acc
    i+=1

acc_final = acc_final/20

print("Average accuracy %2.3f" % acc_final)
```

```
# # Question 2
```

```
#
```

```
# ## Import Data
```

```
# In[9]:
```

```
# Code to import "heart.csv"
```

```
import pandas as pd
data = pd.read_csv('heart.csv')
```

```
# ## Part 1
```

```
# "ST_slope" column should be label encoded because the result has direction, the
rest are one hot encoded
```

```
# In[10]:
```

```
# Code for any variable transformations/pre-processing here
data['ST_Slope'] = data['ST_Slope'].astype('category')
data['ST_Slope'] = data['ST_Slope'].cat.codes
```

```

data_clean = pd.get_dummies(data)

# ## Part 2

# In[11]:

# AUC and accuracies of Gaussian Naive Bayes across 10 random train/test splits
# Make sure to get a table of all AUC and accuracies

from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
i = 0
acc = []
auc = []
while i <= 9:

    df_train, df_test = train_test_split(data_clean, test_size=0.25)
    from sklearn.metrics import accuracy_score

    d=
['Age', 'RestingBP', 'Cholesterol', 'FastingBS', 'MaxHR', 'Oldpeak', 'ST_Slope', 'Sex_F', '
Sex_M', 'ChestPainType_ASY', 'ChestPainType_ATA', 'ChestPainType_NAP', 'ChestPainType_T
A', 'RestingECG_LVH', 'RestingECG_Normal', 'RestingECG_ST', 'ExerciseAngina_N', 'Exercis
eAngina_Y']
    X_train = df_train[d].to_numpy()
    y_train = df_train["HeartDisease"].to_numpy()

    X_test = df_test[d].to_numpy()
    y_test = df_test["HeartDisease"].to_numpy()

    gnb_pend_model = GaussianNB()
    gnb_pend_model.fit(X_train, y_train)

    gnb_pend_pred = gnb_pend_model.predict(X_test)
    gnb_pend_acc = accuracy_score(y_test, gnb_pend_pred)
    gau_roc_auc = roc_auc_score(y_test, gnb_pend_model.predict(X_test))
    acc.append(gnb_pend_acc)
    auc.append(gau_roc_auc)
    i+=1

from pandas import DataFrame
data = {'accuracy':acc,
        'auc':auc}
df = DataFrame(data)
print(df)

# ## Part 3

# When we don't know some of these variables, we can indirectly infer
probabilities by predicting others

```

```
# In[12]:
```

```
# Implement your suggested method here.
```

```
cols=[ 'Age', 'RestingBP', 'FastingBS', 'MaxHR', 'Oldpeak', 'Sex_F', 'Sex_M', 'ChestPainType_ASY', 'ChestPainType_NAP', 'ChestPainType_TA', 'ExerciseAngina_N', 'ExerciseAngina_Y', 'ST_Slope']
```

```
X_train = df_train[cols].to_numpy()
```

```
y_train = df_train["HeartDisease"].to_numpy()
```

```
X_test = df_test[cols].to_numpy()
```

```
y_test = df_test["HeartDisease"].to_numpy()
```

```
gnb_pend_model = GaussianNB()
```

```
gnb_pend_model.fit(X_train, y_train)
```

```
gnb_pend_pred = gnb_pend_model.predict(X_test)
```

```
gnb_pend_acc = accuracy_score(y_test, gnb_pend_pred)
```

```
print("The accuracy using num_of_pendown* features is: ", gnb_pend_acc)
```

```
gau_roc_auc = roc_auc_score(y_test, gnb_pend_model.predict(X_test))
```

```
print('The auc:', gau_roc_auc)
```

```
# ## Part 4
```

```
# In[13]:
```

```
# AUC and accuracies of KNN and WNN using different values of K
```

```
# Make sure to get a table of all AUC and accuracies
```

```
cols=
```

```
['Age', 'RestingBP', 'FastingBS', 'MaxHR', 'Oldpeak', 'Sex_F', 'Sex_M', 'ChestPainType_ASY', 'ChestPainType_NAP', 'ChestPainType_TA', 'ExerciseAngina_N', 'ExerciseAngina_Y', 'ST_Slope']
```

```
df_train, df_test = train_test_split(data_clean, test_size=0.25)
```

```
X_train = df_train[cols].to_numpy()
```

```
y_train = df_train["HeartDisease"].to_numpy()
```

```
X_test = df_test[cols].to_numpy()
```

```
y_test = df_test["HeartDisease"].to_numpy()
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.metrics import accuracy_score
```

```
from sklearn.metrics import roc_auc_score
```

```
acc_knn = []
```

```
auc_knn = []
```

```
acc_wnn = []
```

```
auc_wnn = []
```

```
i= 1
```

```
while i <= 30:
```

```
    clf_1 = KNeighborsClassifier(n_neighbors=i, weights = 'uniform')
```

```
    clf_1 = clf_1.fit(X_train,y_train)
```

```
    clf_1 = clf_1.predict(X_test)
```

```
    clf_acc_1 = accuracy_score(y_test, clf_1)
```

```

clf_auc_1 = roc_auc_score(y_test, clf_1)
acc_knn.append(clf_acc_1)
auc_knn.append(clf_auc_1)
clf_2 = KNeighborsClassifier(n_neighbors=i, weights = 'distance')
clf_2 = clf_2.fit(X_train, y_train)
clf_2 = clf_2.predict(X_test)
clf_acc_2 = accuracy_score(y_test, clf_2)
clf_auc_2 = roc_auc_score(y_test, clf_2)
acc_wnn.append(clf_acc_2)
auc_wnn.append(clf_auc_2)
i+=1

```

```

from pandas import DataFrame
data = {'k':
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30],
        'Accuracy (KNN)': acc_knn,
        'Accuracy (WNN)': acc_wnn,
        'AUC (KNN)': auc_knn,
        'AUC (WNN)': auc_wnn}
df = DataFrame(data)
print(df)

```

# In[ ]:

# In[ ]:

# In[ ]: